



# Sequential Solutions in Machine Scheduling Games

Cong Chen<sup>1</sup>, Paul Giessler<sup>2</sup>, Akaki Mamageishvili<sup>3(✉)</sup>, Matúš Mihalák<sup>2</sup>,  
and Paolo Penna<sup>4</sup>

<sup>1</sup> School of Business Administration, South China University of Technology,  
Guangzhou, China

<sup>2</sup> Department of Data Science and Knowledge Engineering, Maastricht University,  
Maastricht, The Netherlands

<sup>3</sup> Department of Management, Technology and Economics, ETH Zurich,  
Zurich, Switzerland  
[amamageishvili@ethz.ch](mailto:amamageishvili@ethz.ch)

<sup>4</sup> Department of Computer Science, ETH Zurich, Zurich, Switzerland

**Abstract.** We consider the classical machine scheduling, where  $n$  jobs need to be scheduled on  $m$  machines, and where job  $j$  scheduled on machine  $i$  contributes  $p_{i,j} \in \mathbb{R}$  to the load of machine  $i$ , with the goal of minimizing the makespan, i.e., the maximum load of any machine in the schedule. We study inefficiency of schedules that are obtained when jobs arrive sequentially one by one, and the jobs choose themselves the machine on which they will be scheduled, aiming at being scheduled on a machine with small load. We measure the inefficiency of a schedule as the ratio of the makespan obtained in the worst-case equilibrium schedule, and of the optimum makespan. This ratio is known as the *sequential price of anarchy* (**SPoA**). We also introduce two alternative inefficiency measures, which allow for a favorable choice of the order in which the jobs make their decisions. As our first result, we disprove the conjecture of [22] claiming that the sequential price of anarchy for  $m = 2$  machines is at most 3. We show that the sequential price of anarchy grows at least linearly with the number  $n$  of players, assuming arbitrary tie-breaking rules. That is, we show  $\mathbf{SPoA} \in \Omega(n)$ . Complementing this result, we show that  $\mathbf{SPoA} \in O(n)$ , reducing previously known exponential bound for 2 machines. Furthermore, we show that there exists an order of the jobs, resulting in makespan that is at most linearly larger than the optimum makespan. To the end, we show that if an authority can change the order of the jobs adaptively to the decisions made by the jobs so far (but cannot influence the decisions of the jobs), then there exists an adaptive ordering in which the jobs end up in an optimum schedule.

**Keywords:** Machine scheduling · Price of anarchy · Price of stability

## 1 Introduction

We consider the classical optimization problem of scheduling  $n$  jobs on  $m$  *unrelated* machines. In this problem, each job has a (possibly different) processing

time on each of the  $m$  machines, and a schedule is simply an assignment of jobs to machines. For any such schedule, the load of a machine is the sum of all processing times of the jobs assigned to that machine. In this optimization problem, the objective is to find a schedule minimizing the *makespan*, that is, the maximum load among the machines.

In the *game-theoretic* version of this scheduling problem, also known as the *load balancing game*, jobs correspond to players who *selfishly* choose the machine to which the job is assigned. The cost of a player is the *load* of the machine to which the player assigned its own job. Such a setting models, for example, the situation where the machines correspond to servers, and the communication with a server has a latency that depends on the total traffic to the server.

The decisions of the players lead to some *equilibrium* in which no player has an incentive to deviate, though the resulting schedule may not necessarily be optimal in terms of makespan. Such an equilibrium might have a rather high *social cost*, that is, the makespan of the corresponding schedule<sup>1</sup> is not guaranteed to be the optimal one, as in Example 1 below.

*Example 1 (two jobs on two unrelated machines [5]).* Consider two jobs and two unrelated machines, where the processing times are given by the following table:

	job 1	job 2
machine 1	1	$\ell$
machine 2	$\ell$	1

The allocation represented by the gray box is a pure Nash equilibrium in the load balancing game (if a job moves to the other machine, its own cost increases from  $\ell$  to  $\ell + 1$ ), and has makespan  $\ell$ . The optimal makespan is 1 (swap the allocations). This example shows that the makespan of an equilibrium can be arbitrarily larger than the optimum.

The inefficiency of equilibria is a central concept in algorithmic game theory. Typically, one aims to quantify the *efficiency loss* resulting from a *selfish behavior* of the players, where the loss is measured in terms of the social cost. Arguably, the two most popular measures of inefficiency of equilibria are the *price of anarchy* (**PoA**) [27] and the *price of stability* (**PoS**) [4], which, intuitively, consider the *most pessimistic* and the *most optimistic* scenario:

- The *price of anarchy* is the ratio of the cost of the *worst* equilibrium over the *optimal social cost*;
- The *price of stability* is the ratio of the cost of the *best* equilibrium over the *optimal social cost*.

<sup>1</sup> When each player chooses deterministically one machine, this definition is obvious. When equilibria are *mixed* or randomized, each player chooses one machine according to some probability distribution, and the social cost is the expected makespan of the resulting schedule.

The price of anarchy corresponds to the situation (anarchy) in which there is no authority, and players converge to some equilibrium by themselves. In the price of stability, one envisions that there are means to suggest the players how to play, and if that is an equilibrium, then they will indeed follow the suggestion, as no unilateral deviation can improve a player's individual cost. Furthermore, the price of stability provides a lower bound on the efficiency loss of an equilibrium outcome, if, for example, no equilibrium is actually a social optimum.

Example 1 thus shows that the **price of anarchy** of load balancing games is **unbounded even for two jobs and two machines**. Interestingly, the price of stability instead is one ( $\text{PoS} = 1$ ), for any number of jobs and any number of machines. This is because there is always an optimal solution that is also a pure Nash equilibrium [17] (see Sect. 1.3 for details). In a pure Nash equilibrium, players choose their strategies deterministically, as opposed to *mixed* Nash equilibria. In this work, we will also focus on the case in which players act deterministically, though in a sequential fashion (see below).

As the price of anarchy for unrelated machines is very high (unbounded in general), one may ask whether Nash equilibria are really what happens as an outcome in the game, or whether a central authority, which cannot influence the choices of the players (jobs), may alter some aspects of the scheduling setting, and as a result, improve the performance of the resulting equilibria.

Motivated by these issues, in [28] the authors consider the variant in which players, instead of choosing their strategies simultaneously, play sequentially taking their decisions based on the previous choices and also knowing the order of players that will make play. Formally, this corresponds to an *extensive-form game*, and the corresponding equilibrium concept is called a *subgame-perfect equilibrium*. Players always choose their strategy deterministically. The resulting inefficiency measure is called the *sequential price of anarchy* (**SPoA**).

There are two main motivations to study a sequential variant of the load balancing game. First, assuming that all players decide simultaneously to choose the machine to process their jobs is a too strong and unnatural modeling assumption in many situations; furthermore, expecting that all players choose the worst-case machine, as was the case in Example 1, is unnatural as well. Second, one may have the power to explicitly ask the players to make sequential decisions, and make this the policy, which the players are aware of, with the view of lowering the loss of efficiency of the resulting equilibrium schedules. In a sense, such an approach of adjusting the way the players access the machines resembles *coordination mechanisms* [11], which are scheduling policies aiming to achieve a small price of anarchy (see Sect. 1.3 for more details).

## 1.1 Prior Results (SPoA for Unrelated Machines)

The first bounds on the **SPoA** for unrelated machines have been obtained in [28], showing that

$$n \leq \text{SPoA} \leq m \cdot 2^n.$$

Therefore, **SPoA** is *constant* for a constant number of machines and jobs, while **PoA** is *unbounded* even for two jobs and two machines (recall Example 1).

The large gap in the previous bound naturally suggests the question of what happens for *many jobs* and *many machines*. This was addressed by [7] which improved significantly the prior bounds by showing that

$$2^{\Omega(\sqrt{n})} \leq \mathbf{SPoA} \leq 2^n.$$

At this point, one should note that these lower bounds use a *non-constant* number of machines. In other words, it still might be possible that for a *constant number of machines* the  $\mathbf{SPoA}$  is constant. For *two machines*, [22] proved a lower bound  $\mathbf{SPoA} \geq 3$ , and in the same work the authors made the following conjecture:

*Conjecture 1* [22]. *For two unrelated machines,  $\mathbf{SPoA} = 3$  for any number of jobs.*

## 1.2 Our Contributions

In this paper, we disprove Conjecture 1 by showing that in fact,  $\mathbf{SPoA}$  on two machines is *not even constant*. Indeed, it must grow linearly and the conjecture fails already for few jobs:

- For *five jobs* we have  $\mathbf{SPoA} \geq 4$  (Theorem 2);
- In general, with arbitrary tie-breaking rules, it holds that  $\mathbf{SPoA} \geq \Omega(n)$  (Theorem 3).

Note that the result of Theorem 3 uses suitable player-specific tie-breaking rules (see Definition 1). We discuss the implications of using tie-breaking rules more in detail at the end of this subsection.

While Theorem 2 settles the conjecture, the result of Theorem 3 says that  $\mathbf{SPoA}$  is non-constant already for two machines (as the number of jobs grows) for generic tie-breaking rules. We actually conjecture that there exist instances for which the  $\mathbf{SPoA}$  is unbounded without having ties. Moreover, it implies a *strong separation* with the case of *identical* machines, where  $\mathbf{SPoA} \leq 2 - \frac{1}{m}$ , for any number  $m$  of machines [22]. In Theorem 4 we show that  $\mathbf{SPoA}$  is upper bounded by  $2(n - 1)$ , reducing the exponential upper bound obtained in [7] for arbitrarily many machines to linear bound for 2 machines.

The original idea behind the notion of price of stability ( $\mathbf{PoS}$ ) is that an authority can suggest to the players how to play:

*[...] The best Nash equilibrium solution has a natural meaning of stability in this context – it is the optimal solution that can be proposed from which no user will defect. [...] As a result, the global performance of the system may not be as good as in a case where a central authority can simply dictate a solution; rather, we need to understand the quality of solutions that are consistent with self-interested behavior* [4].

We borrow this idea of an authority suggesting desirable equilibria. Specifically for our setting, the authority suggests the order in which players make their

decisions, so to induce a good equilibrium. This can be viewed as the price of stability (**PoS**) for these sequential games. We introduce this notion in two variants (a weaker and a stronger):

- *Sequential Price of Stability (SPoS)*. The authority can choose the order of the players' moves. This order determines the tree structure of the corresponding game.
- *Adaptive Sequential Price of Stability (adaptive SPoS)*. The authority decides the order of the players' moves *adaptively* according to the choices made at each step.

The study of these two notions for two unrelated machines is also motivated by our lower bound, and by the lack of any good upper bound on this problem. We prove the following upper bounds for two unrelated machines (Theorems 5 and 6):

$$\mathbf{SPoS} \leq \frac{n}{2} + 1, \quad \text{adaptive SPoS} = 1.$$

The next natural question is to consider *three* or more machines. Here we show an impossibility result, namely **adaptive SPoS**  $\geq 3/2$  already for three machines (Theorem 7). That is, even with the strongest type of adaptive authority, it is not possible to achieve the optimum. This shows a possible disadvantage of having players capable of complex reasoning, like in extensive-form games. In the classical strategic-games setting, where we consider pure Nash equilibrium, here is an optimum which *is* an equilibrium, that is, **PoS** = 1 for any number of machines and jobs. This result follows from [17] (see Sect. 1.3 for details).

As mentioned above, some of our results rely on the use of a suitable tie-breaking rules. Using tie-breaking rules to prove lower bounds on the **SPoS** is not new: in [25] the authors showed that, in *routing games*, the sequential price of anarchy is *unbounded*. Their proof is based on carefully chosen tie-breaking rules. This way of using tie-breaking rules is not part of the players' strategy interactions. In contrast, some works consider settings where among equivalent choices, each player  $i$  can use the one that hurts prior agents who chose a strategy that player  $i$  would prefer they had not chosen (see [30]).

### 1.3 Further Related Work

The load balancing games considered in this work are one of the most studied models in algorithmic game theory (see, e.g., [2, 15, 16, 18, 19, 26, 27]). In all these works, players correspond to jobs, their cost is the load of the machine they choose, and the social cost is defined as the makespan of the jobs allocation. In particular, the seminal paper [27] which introduced the concept of the price of anarchy, considers the case of *identical* and *related* machines, two simpler versions of unrelated machines (related machines is the setting where each machine has a speed, each job has a certain size, and the processing time equals the job size divided by the machine speed; the case of identical machines is the restriction in which all speeds are the same).

Interestingly, the price of anarchy for *related* or *identical* machines is much better than in the case of unrelated machines (where the price of anarchy is unbounded). Indeed, for related and identical machines, the price of anarchy is *bounded* for any *constant number of machines* [15, 16, 18–20, 26, 27] (some of these results give bounds also for *mixed* Nash equilibria). Specifically, for pure Nash equilibria,  $\mathbf{PoA} = (2 - \frac{2}{m+1})$  for identical machines as implied by the analysis of [20], while  $\mathbf{PoA} = O(\frac{\log m}{\log \log m})$  for related machines [15].

As already mentioned above, the  $\mathbf{PoS}$  for *unrelated* machines is 1. This is due to the work [17] which shows that, starting from any schedule, an iterative process of applying unilateral improving-strategy changes of players leads to a pure Nash equilibrium (the same property has been observed earlier in [21] for related machines). This condition implies the existence of a pure Nash equilibrium.

Requiring that players make their decisions sequentially, according to a given and known order can be seen as a mean of a central authority that can control access to the resources (machines), but not the choices of the players (jobs). In this sense, changing the access from simultaneous to sequential can be seen as a kind of control mechanism like a *coordination mechanism* [11]. In load balancing games where the cost of a player (job) is the completion time of the job (and not the total load of the machine on which the job is scheduled), a coordination mechanism is a scheduling policy, one for every machine, which determines the order of the jobs in which they will be scheduled on the machine. The scheduling policy needs to be fixed and (publicly) known to the players. For load balancing games in normal form (i.e., where players make simultaneous decisions, as opposed to the sequential decisions, which we consider in this paper), coordination mechanisms have been studied both for the version where the social cost is the makespan (see, e.g., [6, 8, 9, 24] and the references therein), or the total (weighted) completion time (see, e.g., [1, 12, 14, 23, 31] and the references therein).

As already discussed above, the concept of a sequential price of anarchy is not new. In addition to the results for unrelated machines discussed in Sect. 1.1, the sequential price of anarchy has been studied also for other games. These include congestion games with affine delay functions [25], isolation games [3], and network congestion games [13]. Interestingly, the latter work shows that the sequential price of anarchy for these games is *unbounded*, as opposed to the price of anarchy which was known to be  $5/2$ .

Naturally, there is a huge literature on the classical algorithm-theoretic research on machine scheduling, see, e.g., the textbook [29] and the survey [10] for fundamental results and further references.

## 2 Preliminaries

In unrelated machine scheduling there are  $n$  jobs and  $m$  machines, and the processing time of job  $j$  on machine  $i$  is denoted by  $p_{ij}$ . A solution (or schedule) consists of an assignment of each job to one of the machines, that is, a vector  $s = (s_1, \dots, s_n)$  where  $s_j$  is the machine to which job  $j$  is assigned to. The *load*

$l_i(s)$  of a machine  $i$  in schedule  $s$  is the sum of the processing times of all jobs allocated to it, that is,  $l_i(s) = \sum_{j:s_j=i} p_{ij}$ . The social cost of a solution  $s$  is the *makespan*, that is, the maximum load among all machines.

Each job  $j$  is a *player* who attempts to minimize her own cost  $cost_j(s)$ , that is, the load of the machine she chooses:  $cost_j(s) = l_{s_j}$ . Every player  $j$  decides  $s_j$ , the assignment of job  $j$  to a machine. The combination of all players strategies gives a schedule  $s = (s_1, \dots, s_n)$ .

In the extensive-form version of these games, players play sequentially; they decide their strategies based on the choices of the previous players and knowing that the remaining players will play rationally. We consider a *full information* game. As players enter the game sequentially, they can compute their optimal moves by the so-called *backward induction*: the last player makes her move greedily, the player before the last makes the move also greedily (taking into account what the last player will do), and so on. Any game of this type can be modeled by a *decision tree*, which is a rooted tree where the non-leaf vertices correspond to the players in certain states, while edges correspond to the strategies available to the players in a given state.

Each leaf corresponds to a solution (schedule), which is simply the strategies on the unique leaf-to-root path. Given the processing times  $P = (p_{ij})$ , the players can compute the loads on the machines in each of the leaves. In case of ties, all players know the deterministic tie-breaking rules of all the other players. A player can calculate what the final outcome would be for each of her strategies, and choose the strategy that minimizes her cost. This method is called backward induction. Strategies obtained in this way for each internal node constitute what is called the *subgame-perfect equilibrium*: for each subtree, we know what is the outcome achieved by the players in this subtree if they play rationally. We usually represent the strategies (edges) that are chosen by players in the **subgame perfect equilibrium** in **bold**, and the other strategies as *dashed* edges.

It is easy to see that a subgame-perfect equilibrium always exists and it is unique, for given tie-breaking rules. On the other hand, its computation is difficult, as proved in [28]:

**Theorem 1** [28]. *Computing the outcome of a subgame perfect equilibrium in Unrelated Machine Scheduling is PSPACE-complete.*

*Notation and Formal Definitions.* We consider  $n$  jobs and  $m$  machines, denoted by  $J = (J_1, J_2, \dots, J_n)$  and  $M = (M_1, M_2, \dots, M_m)$  respectively. The processing times are given by a matrix  $P = (p_{ij})$ , with  $p_{ij}$  being the processing time of job  $J_j$  on machine  $M_i$ . The set of all such nonnegative  $n \times m$  matrices is denoted by  $\mathcal{P}_{n,m}$  and it represents the possible instances of the game. For any  $P \in \mathcal{P}_{n,m}$  as above, we denote by  $\mathcal{T}_{n,m}$  the set of all possible depth- $n$ , complete  $m$ -ary decision trees where each path from the root to a leaf contains every job (player) exactly once. The whole game (and the resulting subgame perfect equilibrium) is fully specified by  $P$ ,  $T$ , and the *tie-breaking rule* used by the players. The most general – worst case – scenario is that ties are arbitrary (see Definition 1). In the following, we do not specify the dependency on the ties, and simply denote by

$SPE(P, T)$  the cost (makespan) of the subgame perfect equilibrium of the game. One type of worst-case analysis is to assume the players' order to be adversarial, and the tree  $T$  being chosen accordingly. This is the same as saying that players arrive in a fixed order (say  $J_1, J_2, \dots, J_n$ ) and their costs  $P$  is chosen in an adversarial fashion. In this case, we simply write  $SPE(P)$  as the tree structure is fixed. For a fixed order  $\sigma$  (a permutation) of the players, and costs  $P$ , we also write  $SPE(P, \sigma)$  to denote the quantity  $SPE(P, T)$  where  $T$  is the tree resulting from this order  $\sigma$  of the players. The optimal social cost (makespan) is denoted by  $OPT(P)$ .

We next introduce formal definitions to quantify the inefficiency of subgame perfect equilibria in various scenarios (from the most pessimistic to the most optimistic). The *sequential price of anarchy* (**SPoA**) compares the worst subgame perfect equilibrium with the optimal social cost,

$$\mathbf{SPoA} = \sup_{P \in \mathcal{P}_{n,m}} \frac{SPE(P)}{OPT(P)}.$$

In the *sequential price of stability* (**SPoS**), we can choose the order  $\sigma$  in which players play depending on the instance  $P$ . The resulting subgame perfect equilibrium has cost  $SPE(P, \sigma)$ , which is then compared to the optimum,

$$\mathbf{SPoS} = \sup_{P \in \mathcal{P}_{n,m}} \min_{\sigma \in \mathcal{S}_n} \frac{SPE(P, \sigma)}{OPT(P)},$$

where  $\sigma$  ranges over all permutations  $\mathcal{S}_n$  of the  $n$  players. In *adaptive sequential price of stability* (**adaptive SPoS**), we can choose the whole structure of the tree, meaning that for each choice of a player, we can adaptively choose which player will play next. This means that every path from any leaf to the root corresponds to a permutation of the players. The adaptive price of stability is then defined as

$$\mathbf{adaptive SPoS} = \sup_{P \in \mathcal{P}_{n,m}} \min_{T \in \mathcal{T}_{n,m}} \frac{SPE(P, T)}{OPT(P)}.$$

Note that by definition **adaptive SPoS**  $\leq$  **SPoS**  $\leq$  **SPoA**.

### 3 Linear Lower Bound for SPoA

In this section, we consider the sequential price of anarchy for *two* unrelated machines. In [22] the authors proved a lower bound **SPoA**  $\geq 3$  for this case, and they conjectured that this was also a tight bound. We show that unfortunately this is not the case: Already for five jobs, **SPoA**  $\geq 4$ , and with more jobs the lower bound grows linearly, i.e., **SPoA**  $= \Omega(n)$ .

#### 3.1 A Lower Bound for $n = 5$ Players

**Theorem 2.** *For two machines and at least five jobs, the **SPoA** is at least 4.*



### 3.2 Linear Lower Bound

Extending the construction for  $n = 5$  players is non-trivial as this seems to require rather involved constants that multiply the  $\varepsilon$  terms. However, we notice that these terms only help to induce more involved tie-breaking rules of the following form:

**Definition 1 (arbitrary tie-breaking rule).** *We say that the tie-breaking rule is arbitrary if each player uses a tie-breaking rule between machines which possibly depends on the allocation of all players.*

The following theorem gives our general lower bound:

**Theorem 3.** *Even for two machines, the **SPoA** is at least linear in the number  $n$  of jobs, in the case of arbitrary tie-breaking rule.*

Note that it is important to have seemingly equivalent jobs  $J_2$  and  $J_3$ . They use different tie-breaking rules, which creates the asymmetry between them and increases the **SPoA**.

We solved linear programs with strict inequalities obtained from the subgame perfect equilibria tree structure given in the example from the proof of Theorem 3, by introducing small  $\varepsilon$  for strict inequalities. We found solutions for  $n = 8$  and  $n = 11$ , that is linear programs are feasible. Therefore, at least for small  $n$ 's we can drop the assumption about tie-breaking rules. As the solutions replace the  $\varepsilon$  terms by rather more complicated coefficients, we do not present them here. For the general case, we conjecture that the statement of Theorem 3 holds without the assumption on the tie-breaking rules, and that the latter are merely used to make the analysis easier:

*Conjecture 2.* For two machines, the **SPoA** is at least linear in  $n$ .

## 4 Linear Upper Bound for SPoA

*Additional Notation.* To prove the upper bound for **SPoA**, we introduce some additional notation. We define a vector  $D = (d_1, d_2)$  of initial load on the machines before the jobs play the game. Consequentially, the load of each machine  $i$  becomes

$$l_i(D, s) = d_i + \sum_{j: s_j = i} p_{ij},$$

where  $s = (s_1, s_2, \dots, s_n)$  is the schedule (SPE) achieved by the jobs playing the game with initial load  $D$  on the machines; the cost of each job  $j$  is

$$\text{cost}_j(D, s) = l_{s_j}(D, s).$$

The notation for the makespan is renewed as  $SPE_D(P)$  for the SPE with initial load  $D$ . Additionally, we define  $\Delta SPE(P)$  as the maximum possible increase of the makespan due to the players, with processing time  $P$ , for any initial load  $D$ :

$$\Delta SPE(P) = \sup_D \{SPE_D(P) - \|D\|_\infty\}.$$

Moreover, for a given  $P$ , we use  $P_{[u:v]}$  to represent the processing times only for jobs  $(J_u, J_{u+1}, \dots, J_v)$ , that is,  $P_{[u:v]} = (p_{ij})$  where  $j = u, u+1, \dots, v$ .

We first show a key lemma showing that each job can only contribute a certain amount (bounded by the total minimum processing time) to the makespan:

**Lemma 1.**  $\Delta SPE(P_{[\ell:n]}) - \Delta SPE(P_{[\ell+1:n]}) \leq \sum_{j=\ell}^n \min_i p_{ij}$  for  $\ell = 1, 2, \dots, n-1$ .

**Theorem 4.** For two machines, the **SPoA** is at most  $2(n-1)$ .

*Proof.* Applying Lemma 1, we have

$$\begin{aligned} \Delta SPE(P_{[1:n]}) &\leq \Delta SPE(P_{[2:n]}) + \sum_{j=1}^n \min_i p_{ij} \\ &\leq \Delta SPE(P_{[3:n]}) + 2 \sum_{j=1}^n \min_i p_{ij} \\ &\leq \dots \\ &\leq (n-1) \sum_{j=1}^n \min_i p_{ij}. \end{aligned}$$

Since the optimal cost is at least  $OPT \geq \sum_{j=1}^n \min_i p_{ij}/2$  (for 2 machines), it follows that

$$\mathbf{SPoA} \leq \frac{\Delta SPE(P_{[1:n]})}{OPT} \leq 2(n-1),$$

which completes the proof.  $\square$

## 5 Linear Upper Bound on the SPoS

In this section, we give a *linear upper bound* on the sequential price of stability for two machines (Theorem 5 below). Unlike in the case of the sequential price of anarchy, here we have the freedom to choose the order of the players. Each player can choose *any* tie-breaking rule. Since we consider a full information setting, the tie-breaking rules are also public knowledge.

Though finding the best order can be difficult, we found that a large set of permutations already gives a linear upper bound on **SPoS**. In particular, it is enough that the authority divides the players into *two groups* and puts players in the first group first, followed by the players from the second group. Inside each group players can form *any order*. The main result of this section is the following theorem:

**Theorem 5.** For two machines, the **SPoS** is at most  $\frac{n}{2} + 1$ .

This result cannot be extended to *three* or more machines, because the third machine changes the logic of the proof. In particular, we can no longer assume that the players on the second machine in the optimal assignment can guarantee low costs for themselves by simply staying on that machine. For two machines, we conjecture that actually there is always an order which leads to the optimum:

*Conjecture 3.* For two machines, the **SPoS** is 1.

Though we are not able to prove this conjecture, in the next section, we introduce a more restricted solution concept, and show that in that case the optimum can be achieved.

## 6 Achieving the Optimum: The Adaptive SPoS

In this section, we study the adaptive sequential price of stability. Unlike the previous models, here we assume that there is some authority, which has full control over the order of the players' arrival in the game. It does not only fix the initial complete order, but can also change the order of arrivals depending on the decision that previous players made. On the other hand, the players still have the freedom to choose any action in a given state, each of them aiming at minimizing her own final cost. The players also know the whole decision tree, and thus the way the authority chooses the order. As in the previous section, each player can use *any* tie-breaking rule, and the tie-breaking rules are also known to all players.

This model is the closest instantiation of a general extensive form game compared to the previously studied models in this paper. In this way, the authority has an option to punish players for deviating from the optimal path (path leading to a social optimum) by placing different players after the deviating decisions of the deviating player. As a result, rational players may achieve much better solutions in the end. The following theorem shows that achieving the optimum solution is possible for 2 machines:

**Theorem 6.** *For two machines, the **adaptive SPoS** is 1.*

The previous result cannot be extended to more than 2 machines:

**Theorem 7.** *For three or more machines, the **adaptive SPoS** is at least  $\frac{3}{2}$ .*

*Proof.* Consider the following instance with three machines and three jobs, where the optimum is shown as gray boxes:

	$J_1$	$J_2$	$J_3$
$M_1$	$4 - \varepsilon$	2	2
$M_2$	4	3	3
$M_3$	6	$6 - \varepsilon$	$6 - \varepsilon$

We distinguish two cases for the first player to move (the root of the tree), and show that in neither case the players will implement the optimum:

1. *The First to Move is  $J_1$ .* This player will choose the cheapest machine  $M_1$ , because none will join this machine. Indeed, the second player to move will choose  $M_2$  knowing that the last one will then choose  $M_3$ .

2. *The First to Move is  $J_2$  or  $J_3$ .* This player will choose  $M_2$  and *not*  $M_1$ . Indeed, if the first player to move, say  $J_2$ , chooses  $M_1$ , then either (I) the other two follow also the optimum (which costs 4 to  $J_2$ ) or (II) they choose another solution, whose cost is at least  $6 - \varepsilon$ . In the latter case, we have the lower bound. In case (I), we argue that choosing  $M_2$  is better for  $J_2$ , because no other player will join: for the following players, being both on machine  $M_1$  is already cheaper than being on  $M_2$  with  $J_2$ .

In the first case, given that  $J_1$  is allocated to  $M_1$ , the cheapest solution costs  $6 - \varepsilon$ . In the second case, one among  $J_2$  or  $J_3$  is allocated to  $M_2$ . The best solution, in this case, costs again  $6 - \varepsilon$ . This completes the proof.  $\square$

*Remark 1.* The following example shows that the analysis of Theorem 6 cannot be extended to 3 machines even in the case of identical machines. Assume that we have  $m = 3$  machines, the initial loads on these machines are  $(0, 2, 6)$  and there are 3 jobs left to be assigned with processing times 7, 5 and 5. Note that the constrained optimum here is  $(10, 9, 6)$ , that is the first job with processing time 7 gets assigned to the second machine  $M_2$ , while both jobs with processing times 5 and 5 get assigned to machine  $M_1$ . On the other hand, if any of these players chooses different machine their cost is strictly decreasing in the subgame perfect equilibrium solution. We did not find any example showing that **adaptive SPoS**  $> 1$  for more than 2 identical machines, unlike the case of unrelated machines.

## 7 Conclusions

In this paper, we disprove a conjecture from [22] and give a linear lower bound construction for the sequential price of anarchy. On the other hand, we show linear upper bound. For the best sequence of players, we prove a linear upper bound, that is 4 times lower than the upper bound for sequential price of anarchy. Moreover, we prove the existence of a sequential extensive game which gives an optimum solution. One possible direction for future research is to investigate whether the sequential price of stability is 1 for any number of *identical* machines. In this work, we give some evidence that the case of three (or more) machines is different from the case of two machines (see Theorem 7 and Remark 1).

Our linear lower bound on the sequential price of anarchy (Theorem 3) suggests that subgame perfect equilibria do not guarantee in the worst case a price of anarchy independent of the number of jobs, even for two machines. Though our lower bound is based on a suitable tie-breaking rule, we believe it holds without any tie being involved (Conjecture 2).

**Acknowledgment.** We are grateful to Thomas Erlebach for spotting a mistake in an earlier proof of Theorem 6 and for suggesting a fix of the proof. We thank Paul Dütting for valuable discussions. We also thank anonymous reviewers and seminar participants of OR 2016 for suggestions that improved the paper.

## References

1. Abed, F., Correa, J.R., Huang, C.-C.: Optimal coordination mechanisms for multi-job scheduling games. In: Schulz, A.S., Wagner, D. (eds.) *ESA 2014*. LNCS, vol. 8737, pp. 13–24. Springer, Heidelberg (2014). [https://doi.org/10.1007/978-3-662-44777-2\\_2](https://doi.org/10.1007/978-3-662-44777-2_2)
2. Andelman, N., Feldman, M., Mansour, Y.: Strong price of anarchy. *Games Econ. Behav.* **2**(65), 289–317 (2009)
3. Angelucci, A., Bilò, V., Flammini, M., Moscardelli, L.: On the sequential price of anarchy of isolation games. *J. Comb. Optim.* **29**(1), 165–181 (2013). <https://doi.org/10.1007/s10878-013-9694-9>
4. Anshelevich, E., Dasgupta, A., Kleinberg, J.M., Tardos, É., Wexler, T., Roughgarden, T.: The price of stability for network design with fair cost allocation. *SIAM J. Comput.* **38**(4), 1602–1623 (2008)
5. Awerbuch, B., Azar, Y., Richter, Y., Tsur, D.: Tradeoffs in worst-case equilibria. *Theoret. Comput. Sci.* **361**(2), 200–209 (2006)
6. Azar, Y., Fleischer, L., Jain, K., Mirrokni, V., Svitkina, Z.: Optimal coordination mechanisms for unrelated machine scheduling. *Oper. Res.* **63**(3), 489–500 (2015)
7. Bilò, V., Flammini, M., Monaco, G., Moscardelli, L.: Some anomalies of farsighted strategic behavior. *Theory Comput. Syst.* **56**(1), 156–180 (2015). [https://doi.org/10.1007/978-3-642-38016-7\\_19](https://doi.org/10.1007/978-3-642-38016-7_19)
8. Caragiannis, I.: Efficient coordination mechanisms for unrelated machine scheduling. *Algorithmica* **66**(3), 512–540 (2013)
9. Caragiannis, I., Fanelli, A.: An almost ideal coordination mechanism for unrelated machine scheduling. *Theory Comput. Syst.* **63**(1), 114–127 (2018). <https://doi.org/10.1007/s00224-018-9857-2>
10. Chen, B., Potts, C.N., Woeginger, G.J.: A review of machine scheduling: complexity, algorithms and approximability. In: Du, D.Z., Pardalos, P.M. (eds.) *Handbook of Combinatorial Optimization*. Springer, Boston (1998). [https://doi.org/10.1007/978-1-4613-0303-9\\_25](https://doi.org/10.1007/978-1-4613-0303-9_25)
11. Christodoulou, G., Koutsoupias, E., Nanavati, A.: Coordination mechanisms. *Theor. Comput. Sci.* **410**(36), 3327–3336 (2009). *Graphs, Games and Computation: Dedicated to Professor Burkhard Monien on the Occasion of his 65th Birthday*
12. Cole, R., Correa, J.R., Gkatzelis, V., Mirrokni, V., Olver, N.: Decentralized utilitarian mechanisms for scheduling games. *Games Econ. Behav.* **92**, 306–326 (2015)
13. Correa, J.R., de Jong, J., de Keijzer, B., Uetz, M.: The inefficiency of Nash and subgame perfect equilibria for network routing. *Math. Oper. Res.* **44**(4), 1286–1303 (2019)
14. Correa, J.R., Queyranne, M.: Efficiency of equilibria in restricted uniform machine scheduling with total weighted completion time as social cost. *Naval Res. Logist. (NRL)* **59**(5), 384–395 (2012)
15. Czumaj, A., Vöcking, B.: Tight bounds for worst-case equilibria. *ACM Trans. Algorithms* **3**(1), 4:1–4:17 (2007)
16. Epstein, L.: Equilibria for two parallel links: the strong price of anarchy versus the price of anarchy. *Acta Informatica* **47**(7), 375–389 (2010)
17. Even-Dar, E., Kesselman, A., Mansour, Y.: Convergence time to Nash equilibria. In: Baeten, J.C.M., Lenstra, J.K., Parrow, J., Woeginger, G.J. (eds.) *ICALP 2003*. LNCS, vol. 2719, pp. 502–513. Springer, Heidelberg (2003). [https://doi.org/10.1007/3-540-45061-0\\_41](https://doi.org/10.1007/3-540-45061-0_41)

18. Feldmann, R., Gairing, M., Lücking, T., Monien, B., Rode, M.: Nashification and the coordination ratio for a selfish routing game. In: Baeten, J.C.M., Lenstra, J.K., Parrow, J., Woeginger, G.J. (eds.) ICALP 2003. LNCS, vol. 2719, pp. 514–526. Springer, Heidelberg (2003). [https://doi.org/10.1007/3-540-45061-0\\_42](https://doi.org/10.1007/3-540-45061-0_42)
19. Fiat, A., Kaplan, H., Levy, M., Olonetsky, S.: Strong price of anarchy for machine load balancing. In: Arge, L., Cachin, C., Jurdziński, T., Tarlecki, A. (eds.) ICALP 2007. LNCS, vol. 4596, pp. 583–594. Springer, Heidelberg (2007). [https://doi.org/10.1007/978-3-540-73420-8\\_51](https://doi.org/10.1007/978-3-540-73420-8_51)
20. Finn, G., Horowitz, E.: A linear time approximation algorithm for multiprocessor scheduling. BIT Numer. Math. **19**(3), 312–320 (1979). <https://doi.org/10.1007/BF01930985>
21. Fotakis, D., Kontogiannis, S., Koutsoupias, E., Mavronicolas, M., Spirakis, P.: The structure and complexity of Nash equilibria for a selfish routing game. In: Widmayer, P., Eidenbenz, S., Triguero, F., Morales, R., Conejo, R., Hennessy, M. (eds.) ICALP 2002. LNCS, vol. 2380, pp. 123–134. Springer, Heidelberg (2002). [https://doi.org/10.1007/3-540-45465-9\\_12](https://doi.org/10.1007/3-540-45465-9_12)
22. Hassin, R., Yovel, U.: Sequential scheduling on identical machines. Oper. Res. Lett. **43**(5), 530–533 (2015)
23. Hoeksma, R., Uetz, M.: The price of anarchy for Minsum related machine scheduling. In: Solis-Oba, R., Persiano, G. (eds.) WAOA 2011. LNCS, vol. 7164, pp. 261–273. Springer, Heidelberg (2012). [https://doi.org/10.1007/978-3-642-29116-6\\_22](https://doi.org/10.1007/978-3-642-29116-6_22)
24. Immorlica, N., Li, L.E., Mirrokni, V.S., Schulz, A.S.: Coordination mechanisms for selfish scheduling. Theoret. Comput. Sci. **410**(17), 1589–1598 (2009)
25. de Jong, J., Uetz, M.: The sequential price of anarchy for atomic congestion games. In: Liu, T.-Y., Qi, Q., Ye, Y. (eds.) WINE 2014. LNCS, vol. 8877, pp. 429–434. Springer, Cham (2014). [https://doi.org/10.1007/978-3-319-13129-0\\_35](https://doi.org/10.1007/978-3-319-13129-0_35)
26. Koutsoupias, E., Mavronicolas, M., Spirakis, P.: Approximate equilibria and ball fusion. Theory Comput. Syst. **36**(6), 683–693 (2003). <https://doi.org/10.1007/s00224-003-1131-5>
27. Koutsoupias, E., Papadimitriou, C.: Worst-case equilibria. In: Meinel, C., Tison, S. (eds.) STACS 1999. LNCS, vol. 1563, pp. 404–413. Springer, Heidelberg (1999). [https://doi.org/10.1007/3-540-49116-3\\_38](https://doi.org/10.1007/3-540-49116-3_38)
28. Leme, R.P., Syrgkanis, V., Tardos, É.: The curse of simultaneity. In: Proceedings of Innovations in Theoretical Computer Science (ITCS), pp. 60–67 (2012)
29. Pinedo, M., Hadavi, K.: Scheduling: theory, algorithms and systems development. In: Gaul, W., Bachem, A., Habenicht, W., Runge, W., Stahl, W.W. (eds.) Operations Research Proceedings 1991. Operations Research Proceedings 1991, vol 1991. Springer, Heidelberg (2012). [https://doi.org/10.1007/978-3-642-46773-8\\_5](https://doi.org/10.1007/978-3-642-46773-8_5)
30. Tranæs, T.: Tie-breaking in games of perfect information. Games Econ. Behav. **22**(1), 148–161 (1998)
31. Zhang, L., Zhang, Y., Du, D., Bai, Q.: Improved price of anarchy for machine scheduling games with coordination mechanisms. Optim. Lett. **13**(4), 949–959 (2018). <https://doi.org/10.1007/s11590-018-1285-3>